

General Regulations.

- Please hand in your solutions in groups of two (preferably from the same tutorial group).
- Your solutions to theoretical exercises can be either handwritten notes (scanned), or typeset using L^AT_EX. For scanned handwritten notes, please ensure they are legible and not blurry.
- For the practical exercises, always provide the (commented) code as well as the output, and don't forget to explain/interpret the latter.
- Please hand in a **single PDF** that includes both the exported notebook and your solutions to the theoretical exercises. Submit the PDF to the Übungsgruppenverwaltung once per group, making sure to include the names of both group members in the submission.
- You can find all the data in the [GitHub Repository](#).

1 The logistic sigmoid

- (a) Compute and simplify the derivative of the binary logistic sigmoid. (1 pt)
- (b) **Bonus:** Show that the tanh function $\tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$, another function sometimes used as the activation function in neural networks, is simply a scaled and shifted version of the binary logistic sigmoid. (2 bonus pts)
- (c) Consider binary classification: Given two points of class 1, one at (1, 1) and one at (2, 2) as well as two points of class 2, at (1, 2) and (2, 3), find a weight vector \mathbf{w} and bias b such that the activation

$$a = \sigma(\mathbf{w}^T \mathbf{x} + b)$$

separates the two classes. (2 pts)

2 Logistic regression: an LLM lie detector

In this exercise, you will implement a lie detection system for Large Language Models (LLMs) using logistic regression (LR). Your lie detector is trained on a dataset of a few thousand hidden activation vectors from LLaMA-3-8B-Instruct.

The hidden activations (at layer 12 of the transformer model) have dimension 4096 and were generated by feeding true and false statements to the LLM, e.g., “The city of Berlin is in Germany.” or “The city of Berlin is in France.”. For each statement, these 4096-dimensional vectors represent the model’s internal state while processing the statement. Each activation vector is labeled as either true (1) or false (0).

Note: You will learn more about LLMs later in the lecture, but you do not need any additional prior knowledge for this exercise. The experiments you consider in this exercise are directly based on the research of Lennart Bürger (former Heidelberg MSc student who published his work at the last NeurIPS¹).

- (a) Use the four provided datasets to train separate Logistic Regression classifiers. Train the Logistic Regression classifier (**without regularization**) on the training set. Evaluate the model’s performance on the test set. Are the activation vectors of true and false statements linearly separable? You may use the Logistic Regression implementation from sklearn (https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.LogisticRegression.html).

Hint: The default implementation uses L2-regularization. You can also play with the tolerance used as stopping criterion. (2 pts)

¹https://proceedings.neurips.cc/paper_files/paper/2024/hash/f9f54762cbb4fe4dbffdd4f792c31221-Abstract-Conference.html

- (b) Now we will investigate how well an LR classifier trained on one dataset generalizes to other datasets (out-of-distribution generalization). For this, train an LR classifier on the cities dataset (once *with* and once *without* regularization) and evaluate its performance on the other three datasets. The sp_en_trans dataset contains true and false Spanish to English translations, e.g., “The Spanish word ‘uno’ means ‘one.’” (True). The prefix `neg_` at the beginning of a dataset name indicates that the dataset contains only negated statements, which include the word “not”, e.g., “The city of Bhopal is not in India.” (False). What do you observe? Does the LR classifier generalize from one topic? Does it generalize from statements without negation to statements with negation? (2 pts)
- (c) Is it possible to train a lie detector that works well for both affirmative and negated statements? Train an LR classifier on the cities and the neg_cities dataset, and evaluate its performance on the sp_en_trans and neg_sp_en_trans datasets. (1 pt)

3 Log-sum-exp and soft(arg)max

The log-sum-exp and soft(arg)max² functions are defined on vectors $\mathbf{z} \in \mathbb{R}^k$ as

$$\text{lse}(\mathbf{z}; \lambda) = \frac{1}{\lambda} \log \left(\sum_{j=1}^K \exp(\lambda z_j) \right) \quad \text{and} \quad \text{soft(arg)max}(\mathbf{z}; \lambda)_k = \frac{\exp(\lambda z_k)}{\sum_{j=1}^K \exp(\lambda z_j)},$$

with a scalar parameter $\lambda \in \mathbb{R}^+$ and $k = 1, \dots, K$.

- (a) On which subset of the vectors $\mathbf{z}^1 = (1, 2, 3)^T$, $\mathbf{z}^2 = (11, 12, 13)^T$, $\mathbf{z}^3 = (10, 20, 30)^T$ does the soft(arg)max yield identical results? Show in general whether the soft(arg)max is invariant under (i) constant offset and (ii) rescaling of its input. (2 pts)
- (b) Make a 2D contour plot of $\text{lse}((z_1, z_2)^T; \lambda)$ for $\lambda \in \{1, 10, 100\}$ and both z_1 and z_2 in the range $[-1, 1]$. Discuss these plots using a contour diagram of $\max(z_1, z_2)$ over the same range. (2 pts)
- (c) Plot the two components of the soft(arg)max, over the same range and the same choices for λ as in (c), but as images instead of contour plots. Compare this to corresponding plots of the two components of the arg max over the 2D vectors, represented as a 2D one-hot vectors instead of indices:

$$\text{onehot}(\arg \max_i z_i) = \begin{cases} (1, 0)^T & \text{if } z_1 > z_2, \\ (0, 1)^T & \text{else.} \end{cases}$$

(1 pt)

- (d) Prove that the derivative of the lse is the soft(arg)max. (1 pt)

- (e) **Bonus:** Prove that

$$\lim_{\lambda \rightarrow \infty} \text{lse}(\mathbf{z}; \lambda) = \max(\mathbf{z}),$$

for all $\mathbf{z} \in \mathbb{R}^n$. (2 bonus pts)

4 Linear regions of MLPs

In this exercise, you will build and investigate two regression Multi Layer Perceptrons (MLPs) using the pytorch (<https://pytorch.org/>) deep learning library. It is okay to use an AI agent for this exercise, but please make sure you understand the resulting code.

Consider an MLP that takes two-dimensional inputs. Each hidden layer consists of a linear transformation (`torch.nn.Linear`) followed by a ReLU activation function defined as $\text{ReLU}(x) = \max(x, 0)$ (`torch.nn.ReLU`). The final layer is a linear transformation without activation function and should produce one scalar output. Formally, for H hidden layers, we have

$$\mathbf{a}_0 = \mathbf{x}, \quad \mathbf{a}_{i+1} = \text{ReLU}(\mathbf{W}_i \mathbf{a}_i + \mathbf{b}_i) \quad \text{for } i \in \{0, \dots, H-1\}, \quad \mathbf{y} = \mathbf{W}_H \mathbf{a}_H + \mathbf{b}_H.$$

²In the literature known as just “softmax”.

Here, $\mathbf{x} \in \mathbb{R}^2$ is the input, $\mathbf{y} \in \mathbb{R}^1$ the output, \mathbf{a}_i are the activations and $\mathbf{W}_i, \mathbf{b}_i$ the weight matrices and bias vectors of the linear layers (the parameters of the model).

- (a) Implement a shallow model with a single hidden layer with 20 neurons. For a tutorial on how to do this with pytorch, see for example https://pytorch.org/tutorials/recipes/recipes/defining_a_neural_network.html. How many parameters does the model have? (2 pts)
- (b) Pytorch automatically takes care of the random initialization of the model. Compute the output for a dense grid of at least 500 by 500 points in the range $\mathbf{x} \in [-10, 10] \times [-10, 10]$ and visualize it as an image. Repeat for a larger range; How far do you have to zoom out to capture all the structure? (1 pts)
- (c) Use `numpy.gradient` to (approximately) compute the spatial gradient of the network output (as an image, from part (b)) and visualize both of its components as images using ‘prism’ as the colormap. What do you observe? (1 pts)
- (d) Implement a deeper model with four hidden layers with 5 neurons each, and repeat part (b) and (c). Interpret your results and compare to the shallow model. (2 pts)

5 Bonus: Number of linear regions

What is the maximum number of linear regions of an MLP with a two-dimensional input and one hidden layer with n neurons and ReLU activations? Hint: Consider the construction in the lecture, representing each hidden neuron as a line in the input space. (4 bonus pts)